
**Using Generative Art Techniques to Explore the Work of Darrell
Viner**

Thomas Carroll

**Submitted in accordance with the requirements for the degree of
Computer Science with Artificial Intelligence (MEng)**

2020/21

40 credits

The candidate confirms that the following have been submitted.

<As an example>

Items	Format	Recipient(s) and Date
Deliverable 1, 2, 3	Report	SSO (DD/MM/YY)
Participant consent forms	Signed forms in envelop	SSO (DD/MM/YY)
Deliverable 4	Software codes or URL	(DD/MM/YY)
Deliverable 5	User manuals	Client, Supervisor (DD/M-M/YY)

Type of project: _____

The candidate confirms that the work submitted is their own and the appropriate credit has been given where reference has been made to the work of others.

I understand that failure to attribute material which is obtained from another source may be considered as plagiarism.

(Signature of Student) _____

Summary

<Concise statement of the problem you intended to solve and main achievements (no more than one A4 page)>

Acknowledgements

Thanks to Rob Sturman for running ‘The Mathematics of Music’ (MATH2340) in which some of the ideas used in this report were introduced to me.

Note that it is not acceptable to solicit assistance on ‘proof reading’ which is defined as the “the systematic checking and identification of errors in spelling, punctuation, grammar and sentence construction, formatting and layout in the test”; see

<http://www.leeds.ac.uk/gat/documents/policy/Proof-reading-policy.pdf>.

Contents

1	Introduction	2
1.1	Problem Overview	3
1.2	Aim	4
1.3	Objectives	4
1.4	Deliverables	5
1.5	Initial Plan	5
1.6	Risk Mitigation	6
2	Graphics	7
2.1	Anatomy of the Work	8
2.2	Polygons	8
2.3	The Grid	10
2.3.1	Demo	10
2.4	Landscape Generation	11
2.4.1	Noise	11
2.5	Navigation	15
2.5.1	Graphical Prompts	16
2.5.2	Controls	16
3	Music	18
3.1	Navigating Sonically	18
3.1.1	Intervals	18
3.1.2	Navigating Space	19
3.1.3	Limitations	20

CONTENTS	1
3.1.4 Processing Demo	21
3.2 Composition	21
3.3 Synthesis	22
3.3.1 Additive Synthesis in Tone.JS	23
3.4 Sequencing	24
3.5 Chords	25
4 Recall	27
4.1 Saving the State	27
4.2 Drawing and Interaction	28
4.3 Recalling parameters	29
5 Ethics	30
References	31
A External Materials	33
B Audio Demo	34

Chapter 1

Introduction

The purpose of this project is to use generative art techniques to explore the spaces created by the works of the artist *Darrell Viner (1947-2001)*. Viner's work included movement, sound, and light, and though primarily working with sculpture, he produced a series of pen plotter drawings. These drawings have been called pioneering works in the field of computer art [12].

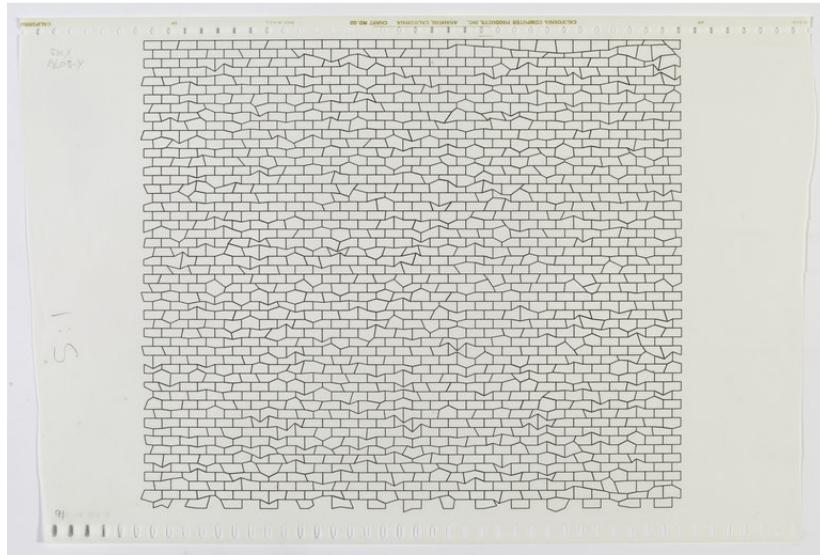


Figure 1.1: *Darrell Viner, Untitled (1974)* ©Victoria and Albert Museum, London.

Generative art is any art that which is created with some system that produces an output, it can either be art which is completed by hand (e.g. a painting) with some system, or stochastic process. Or, more commonly art which is generated by computer program with some initial parameters.

The project will incorporate graphics and sound to create software that can be used to explore the landscapes present in Viner's works. The major problem that needs to be solved in this project is creating an interface that allows the user to interact meaningfully with the program and a set of parameters that might be changed to produce an image. As an extension of the graphics, a musical component will also be produced, this is less about Viner and more about creating a piece of art that stands apart from his work, but is still influenced and adds to the visual aspect.

1.1 Problem Overview

Viner spoke about his work being like a “townscape/landscape”: “Basically it is a self generating program which depends on the start conditions. Thus by altering various aspects of the program, changes in the final image can be achieved. Currently I am after images which have the feel and scale of townscapes/landscapes. The program is modified depending upon whether I consider the images to be working or not: the program has become my personal aesthetic.” [18]

These parameters may for example control the spacing of the grid, the displacement of each point, the number of vertices in a shape, so on. These parameters should be able to be changed to create some sort of ‘landscape’ or ‘topography’ that the user can explore through manipulating the program affecting both the entire grid and single points within it. Later on I will detail methods for users being able to manipulate the parameters and considerations that need to be made with respect to usability.

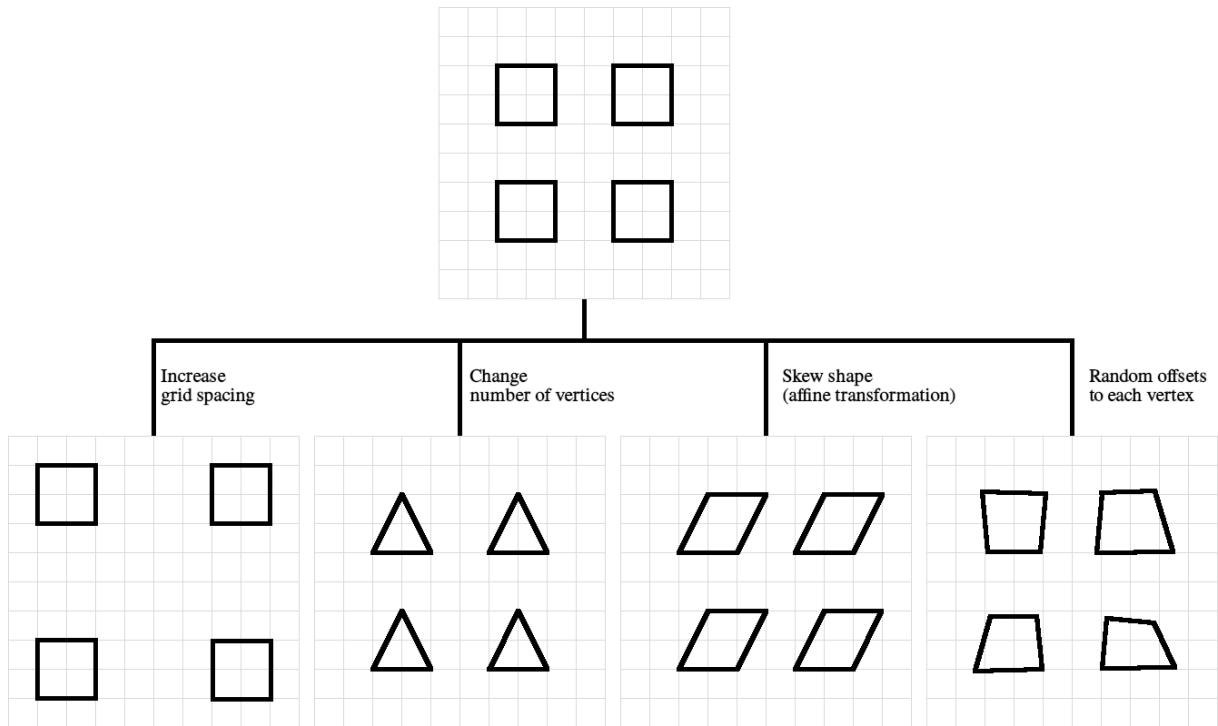


Figure 1.2: Example of changing parameters from an initial state

The program should allow for a user to generate such images and recall them later. Ideally, the program should be more than just a show of the potential configurations of Viner’s work and itself be a unique experience for the user and work to create a dialogue with the original work and contextualise the user’s understanding of his work.

Users of this program may be for example: artists, art scholars, museum guests, museum curators. Each of these have different goals; the artist may be seeking to understand the work and find inspiration in it, the art scholar may be looking for an extended context to the work to

help comprehend the thought process Viner used and relationships between each piece. The museum guest may be at an exhibition (perhaps a virtual one) of Viner's work and have this presented as a piece of work that helps them understand the spaces that are explored by the work. A curator might be evaluating the program for display alongside Viner's work.

1.2 Aim

Using the idea of a virtual museum exhibition of Viner's work: this could be presented alongside a collection of images to show the relationships between each image where in a physical museum this may be able to be explained through the placement of the images in the world, online this is harder to convey with just a gallery on a web page. Therefore the program should allow for a user to both have a casual 'play' with it and to also allow tuning parameters to create images and exploring the images in-depth. Further to aide in the comparative aspect and relationships between images recall and retracing must be possible.

In the context of the art itself there are requirements such as the flexibility of the shape on screen to take on shapes but be ultimately constrained. Since we are replicating a range of possibilities of work of a pen plotter for example, it doesn't make sense for vertices to be very far apart (or even off of the 'page'). This will require development of a set of tools that can be used with the application to calculate point-wise the positions of vertices.

The musical aspect fits in with the idea of an exhibit; a museum version of this would have a pair headphones that the user can wear which would help isolate the piece of art from the rest of the exhibit. In the case of people accessing the piece online the same applies, with audio they can be more immersed in the art. Also, since the audio relates to the graphics on the screen, it can help users navigate the 'landscape' by providing cues.

1.3 Objectives

- To create a program that allows the user to experience the works of Viner through the manipulation of parameters in a way that enables exploration and recall.
- Develop a catalogue of techniques to be used in generative art and music, and to evaluate each for relevance to the project
- Create an interface which allows users to navigate generated spaces
- Create a system such that users are able to recall a state they were in, and find it again
- As an extension of that a way of seeing all previous states, or a map of what the user had seen in a session

1.4 Deliverables

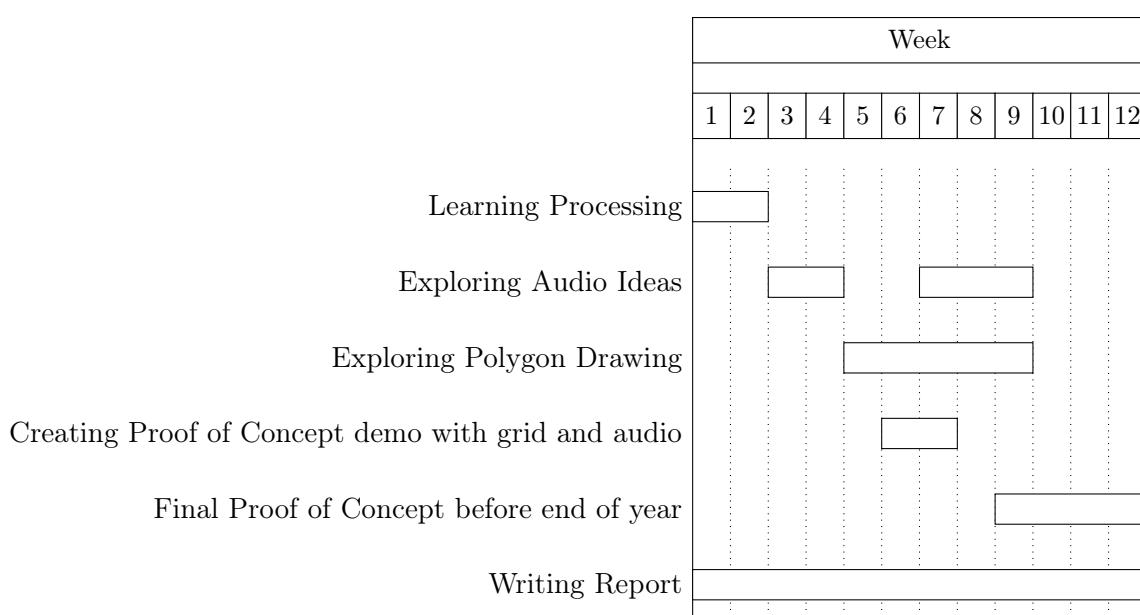
An application written using p5js that should allow the user to explore a landscape, generated to feel like Viner's work. This application should be friendly to use and will include both graphics and sound. The feeling of some sort of topography should be conveyed and the visual should look similar to Viner's work. The audio will help the user navigate the space and correspond to what is on the screen graphically in some way.

There should be another mode within the program that allows users to explore the previous session's movements and recall what they saw. This should be easy to understand and allow for users to see configurations of parameters they like and save them somehow.

1.5 Initial Plan

Being that this project is exploratory in nature, an iterative approach to development makes sense, with small prototypes being created, explored, and built on quickly. As such I have allocated time for ideas to be developed in.

To start I will learn about processing, this library for Java makes it simple to get a framework for graphics, of which the logical can be easily moved to some other framework if need be. Since my tutor's previous work was written in processing, there is already some groundwork complete to build off of. It may be worth exploring other options like p5.js for portability and the ability to embed into webpages for the final implementation, p5.js offers some more mature sound libraries, as well as other web technologies which may be useful for future explorations such as Geo-location and portability to mobile devices.



1.6 Risk Mitigation

Since this project is fairly isolated in nature, there aren't a lot of risks to consider.

If the project had relied on access to the art directly, the COVID-19 pandemic would have affected it, but we mitigated this risk by setting the scope of the project to not include access to the art.

Another risk in this project is scope-creep; since the project is artistic it can be very easy to have many ideas whilst losing the original focus. To combat this I will set specific objectives and deliverables and reach those before exploring any further options.

Chapter 2

Graphics

There are two problems to consider here, how points on a grid are generated, and how they are displayed. To generate points there needs to be a method of procedural generation, i.e. creating an algorithm that can given a set of parameters tell me where a point should or shouldn't be and what other points it should or shouldn't connect to.

The problem of how they are displayed is another issue, how can we convey the idea of moving through a space when the space is simply a grid? The points must move smoothly between configurations to give the illusion of being in a landscape rather than that of a series of static images.

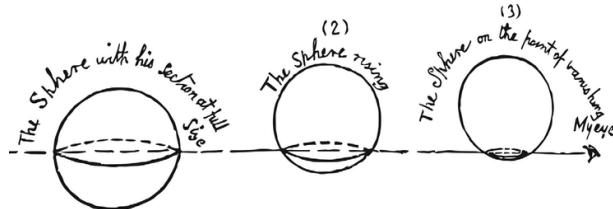


Figure 2.1: Figure reproduced from Flatland [1, p.112]

The landscape is not just 2D, parameters each constitute a dimension, when one is changing we see a 2D slice of the multi-dimensional world as in the book ‘Flatland’ “You cannot indeed see more than one of my sections, or Circles, at a time; for you have no power to raise your eye out of the plane of Flatland; but you can at least see that, as I rise in Space, so my sections become smaller. See now, I will rise; and the effect upon your eye will be that my Circle will become smaller and smaller till it dwindles to a point and finally vanishes.” [1, p.112] A shape (polygon in this instance, sphere in the book) cuts through the plane to show ‘slices’ of itself. Here we are in ‘flatland’ and the parameter space we create allows us to see the work as slices of a higher-dimensional shape (than 2D).

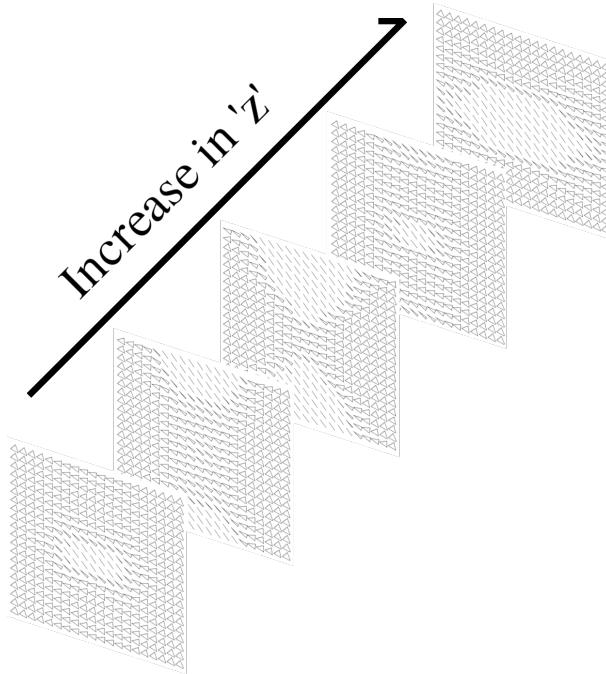


Figure 2.2: Here the ‘z’ parameter increases, each ‘layer’ is transformed between smoothly

2.1 Anatomy of the Work

Viner’s pen plotter work was created by a set of programs, created in **FORTRAN** using a set of subroutines called **PICASO** (PIcture Computer Algorithms SubroutineOriented) [6], this was essentially a line-drawing library, similar to what we are using processing / p5js for. **PICASO**’s use by Viner isn’t well documented but the manual [17] has many subroutines for transforming vertices according to some rules. Some notes on Viner’s work indicate there was definitely mathematical thinking going on in the development.

2.2 Polygons

One aspect the program should be able to do is to grow and shrink polygons between different number of verticies. Viner’s work uses many different number of verticies between images, so we should be able to reproduce this by allowing a smooth transition between polygons.

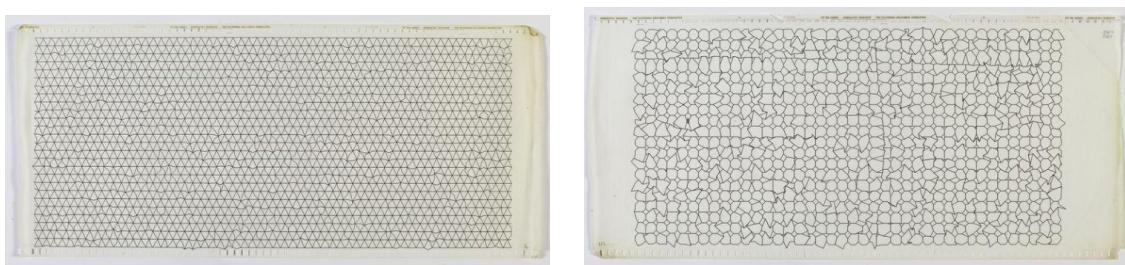


Figure 2.3: Two viner showing the use of polygons

To achieve this a polygon is inscribed on a circle. The points of this polygon of n vertices on the circle are simply given by

$$(x, y) = \left(\cos\left(\frac{2k\pi}{n}\right), \sin\left(\frac{2k\pi}{n}\right) \right) \text{ where } k = 0, 1, \dots, [n]$$

We can note that for integer values of n the rotation will be complete, for non-integer values of n this also works given that we used the floor of n as the limit giving us an incomplete polygon, we can use the fact we have the first coordinate to close the shape. This gives us the smooth transition between integer values of n .

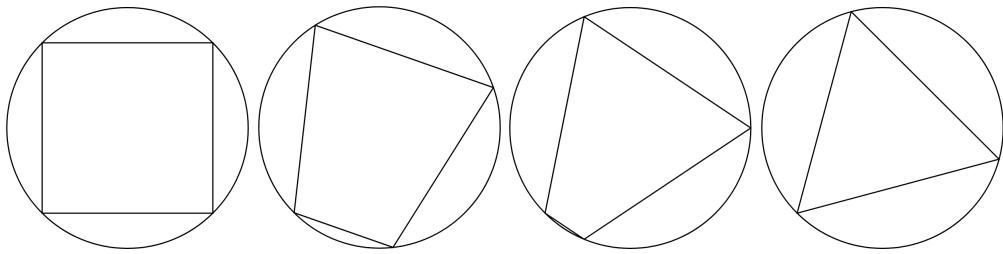


Figure 2.4: Starting at $n = 4$ the shape deforms through $n = 3.5$, $n = 3.2$ and finally $n = 3$

In code we can implement this as such:

```
// center
let x = innerWidth/2;
let y = innerHeight/2;
let n = 4;
let r = 300;

let theta = 3*QUARTER_PI;
dTheta = TWO_PI/n;

beginShape()
  vertex(x + r*cos(theta), y + r*sin(theta));
  for (i = 1; i <= n; i++) {
    theta += dTheta;
    vertex(x + r*cos(theta), y + r*sin(theta));
  }
endShape(CLOSE);
```

An offset to each (x, y) pair can be made, calculated within the for loop, to transform the shape away from regular polyhedra (when n is an integer).

2.3 The Grid

With the motif of the grid being the most obvious visible thing in Viner's art, there needs to be a way to actually draw a grid to the screen; ideally each vertex needs to be able to be separately controlled.

To do this I have adapted John Stell's work on Viner in his workshops [15], using an object-oriented approach; but have created a system where the grid is centred on a given `x,y` coordinate. Essentially there is a screen, and for every vertex at a column and row their relative coordinates need to be calculated. This is simple with the following statement:

```
sx = (x + (gridSize * ((cols/2) - i)));
sy = (y + (gridSize * ((rows/2) - j)));
```

Where `i,j` is the column and row value of the point, `cols,rows` are the total number of columns and rows, and `gridSize` is the pixel size of the grid (which is more of a guide than anything). The grid should also be centred on `x,y` so calling a translate before drawing any points should be done:

```
translate((width/2)+(0-x), (height/2)+(0-y));
```

All of this leads to a system where the centre point has a given coordinate and we can find that coordinate for all other points around it, this also means we can interpolate any parameters from the program to what should be expected at a given coordinate; or we can also generate a terrain away from the main graphics thread and draw them to the screen at a given coordinate.

2.3.1 Demo

Here a grid is prepared with fixed parameters that create distortion around `x=0, y=0`. This point is fixed, and when the program is interacted with the world moves beneath the player rather than the point of distortion being changed, the grid displayed is reflecting the 'terrain' visible from the centre point.

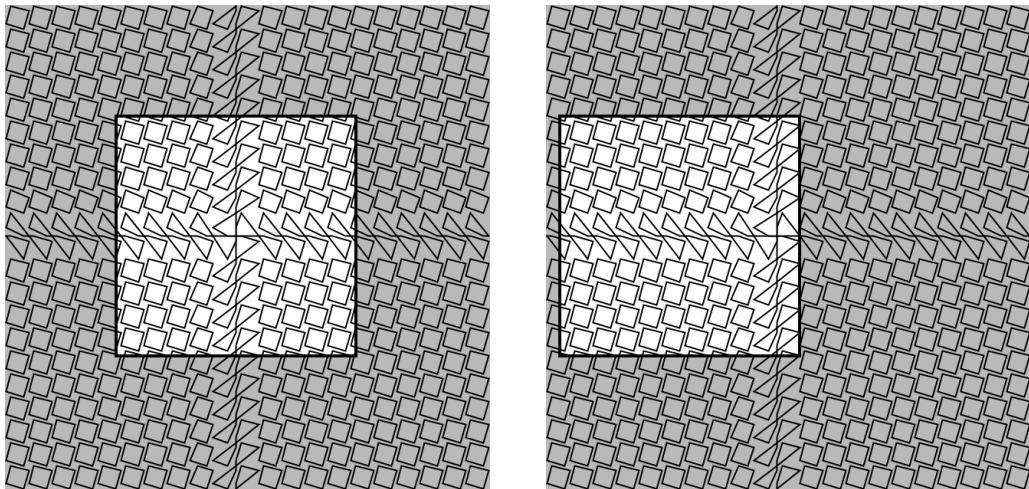


Figure 2.5: The user presses ‘a’ for a short time, decreasing the ‘x’ parameter

2.4 Landscape Generation

The program should generate a ‘landscape’ or ‘topology’, this is an analogy for a space that has an continuous change of the parameters across them. ‘Landscape’ is apt because we’re exploring a set of features generated by the program much like if you were to look at a topographic map. The user should have the feeling of moving through a space like this, more than simply increasing some parameter and immediately seeing a change throughout the grid.

This is the main question of how the program will work on a technical level, how can we create an algorithm or mathematical model that explores the spaces that Viner’s work set out to create?

Dynamical systems may be of interest, and allow for a system to be created where a state evolves into other states following some rules. These states can be deterministic which is important for the objective that we have of recall, but can also be chaotic, which may be aesthetically desirable.

Similarly, fractals may be useful for their self-similarity. Given we’re working with a grid, the ability to have self-similar properties may be considered to be aesthetically useful.

This leads to the choice between having each session using the software be either random in some sense or the same every time. Ideally given a random option to fulfil the ability to recall previous sessions, a seed would be given. It seems then that the random choice contains the static choice and should be the one to be carried out.

2.4.1 Noise

One possible function to consider here is a computer-graphics oriented noise implementation. A desireable property is that of a fairly smooth gradient between extremes, or defined regions in

which values are high.

p5js's built-in noise function is a perlin noise generator. You can pass it up to three coordinates. Perlin noise was designed for computer graphics ¹, and is relatively simple, generating a random grid of vectors and then computing the dot product vectors and their offsets, then finally interpolating to create a more smooth image [10].

p5js also provides a `noiseDetail()` method that provides some control over the 'texture' of the noise. Also for reproducibility `noiseSeed()` allows the programmer to set a seed value for the noise. The `noise()` function takes three coordinate arguments and outputs a number between 0 – 1

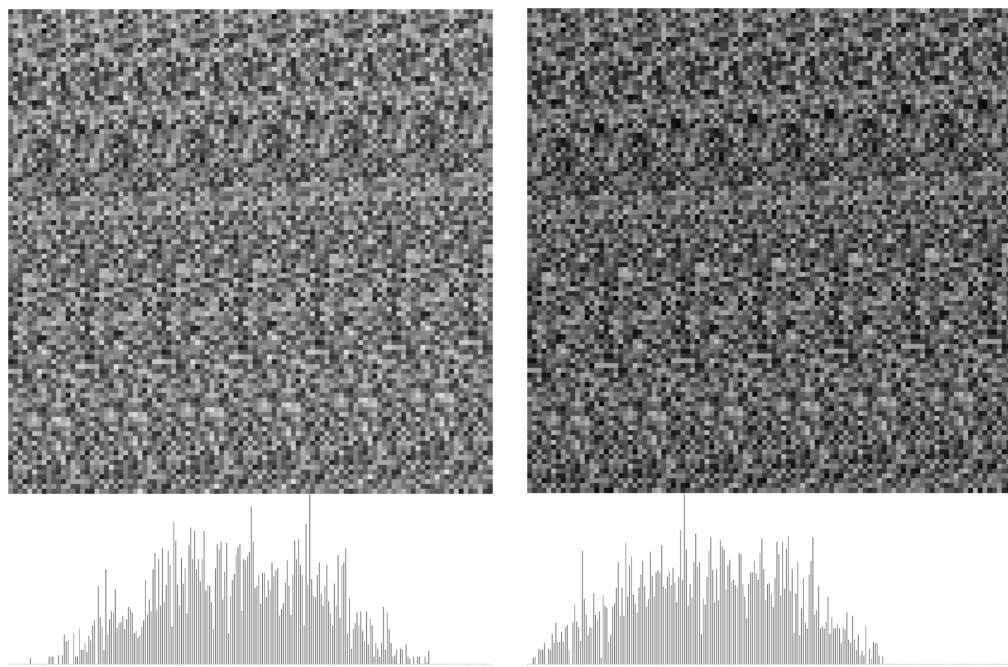


Figure 2.6: Two perlin noise samples (one default the other with `noiseDetail(2, 0.5)` with the same seed, histograms plotted beneath show that the adjustment makes the image overall darker

One option we may wish to be able to have is 'quantising' the noise to some set of values. For example we may wish to have areas of a certain value of n for a polygon. This can be achieved by mapping the noise from 0 – 1 to 0 – n and rounding to the closest integer. This was happening implicitly with the histogram above. We can see that because the distribution is unimodal the most common values will be those towards the center of the range.

¹And won an Oscar for "allow[ing] computer graphics artists to better represent the complexity of natural phenomena" [9]

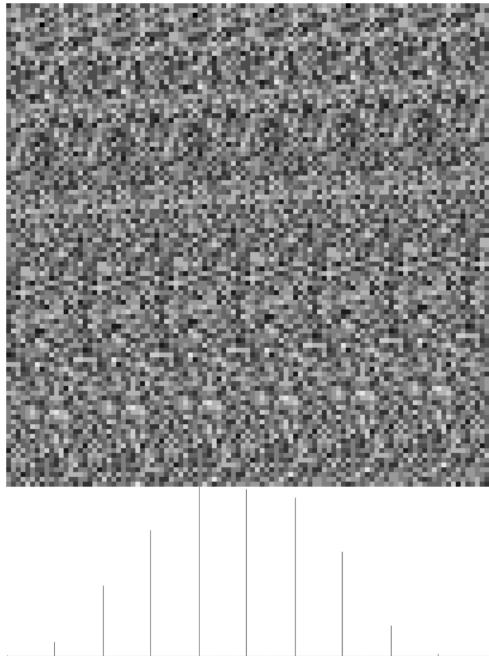


Figure 2.7: This sample contains only ten colours

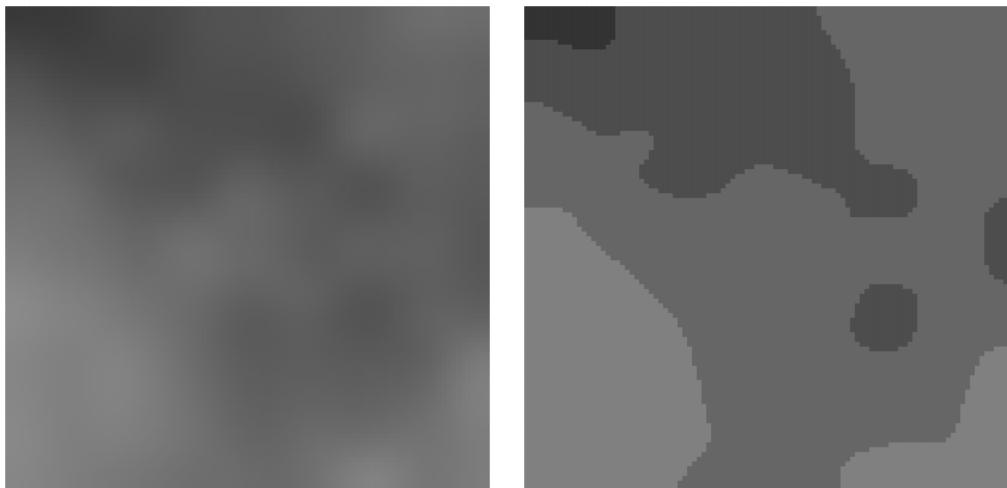


Figure 2.8: Unquantised and quantised noise at 0.001 noiseScale

This effect is much more obvious when the noise scale is a lot smaller, this would then allow for discrete regions between two integer values for a parameter. What if we don't want discrete regions and instead something more like a multimodal distribution i.e. smoothing applied between edges?

We can use a waveform such as the sawtooth to achieve ‘smoothing’ by using it to estimate rounding, adding the value to the wave produces something approximating the ‘staircase’ function of piece-wise rounding. To do this we can use additive synthesis:

$$x + \frac{1}{\pi} \left(- \sum_{k=1}^{\infty} \frac{\sin(2k\pi x)}{k} \right)$$

To use this in code we take only some number of terms, this determines the accuracy of the rounding, the lower the more ‘smooth’, i.e. inexact the output:

```
function approx_round(value, terms) {
    let result = value;

    var innerSum = -sin(2 * PI * value);
    for (i = 2; i <= terms; i++)
        innerSum += (sin((i * 2) * PI * value) / i)

    result += innerSum / PI;

    return result;
}
```

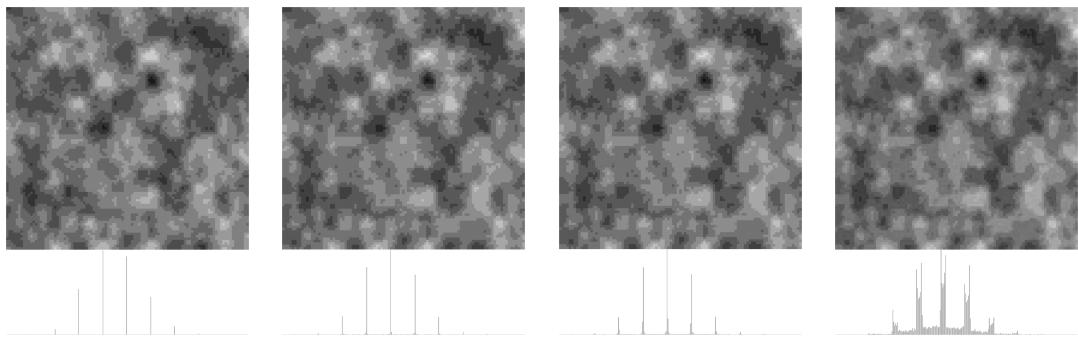


Figure 2.9: Rounded discretely, Terms = 100, Terms = 20, Terms = 1

At such a large scale the difference is hard to perceive but when using the example in Figure 2.8 there’s a clear ‘smoothing’ in the boundaries.



Figure 2.10: Terms = 100, Terms = 10, Terms = 1

Implementing this such that noise controls the number n (verticies in a polygon) we can get results like this (note that the background is shaded darker for lower values of n and lighter for higher values of n)

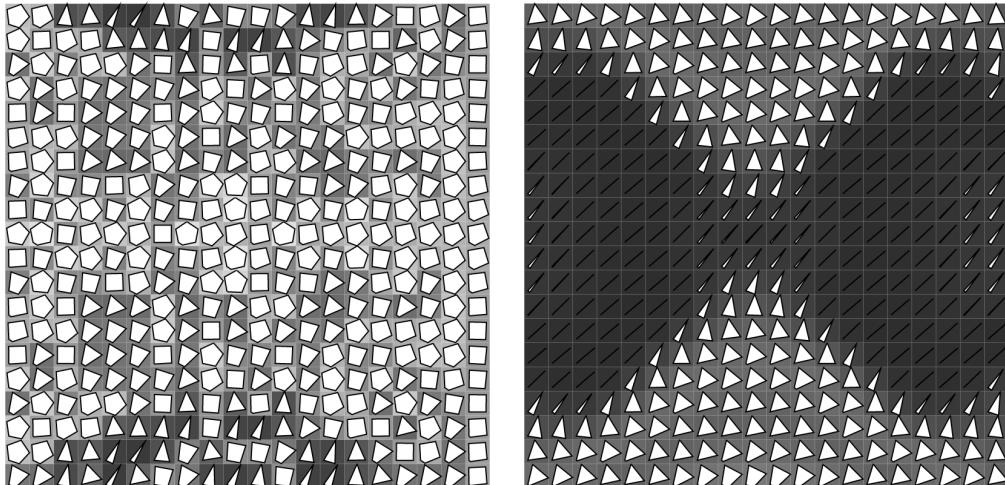


Figure 2.11: Terms = 20, Terms = 1

With user interaction this scheme creates the feeling of moving through regions of these differernt values of n and the parameter for the number of terms allows for adjustments to be made about how ‘hard’ or ‘soft’ the borders between these regions look.

This approach is, I believe somewhat novel, as most techniques for adjusting noise is usually done at a whole-image level. This method is point-wise and thus can work for each polygon or vertex in the grid we have defined. In the field of signal processing the concept of ‘quantisation noise’ is similar but reversed, taking a contiuous signal and analysing the error in the digital discrete signal. Here we are designing a function for an artistic application, and I do not think that it has application to the field of signal processing as it works on discrete to discrete transformations and not continous to discrete.

2.5 Navigation

Navigating higher-dimensional space is of practical interest because a lot of data is higher-dimensional, for example the field of machine learning for example deals with very high dimensional data. Tools like ‘ggobi’ exist to help explore multi-variate data [4] which tend to disregard spacial relations of points in favour of clustering based on shared properties. A novel method is that of the grand tour, through which a series of scatter plots are projected orthogonally into 2D subspaces from the higher-dimensional space and moved then between continuously to describe multi-variate data [2].

Whilst these approaches are useful for data, what we’re dealing with here is generating spaces to explore, these spaces are less about trends in data and more about geometric exploration through many continuous parameters.



Figure 2.12: Mario turns 90° around the z-axis (*Super Paper Mario (2007)* pub. Nintendo)

Video games also have explored higher-dimensional spaces. An interesting concept is that of 2D characters being able to navigate 3D spaces. *Super Paper Mario (2007)* for the Nintendo Wii is an example of this, the player can turn the world 90 degrees about the y-axis to explore the depth of the z-axis and progress through the game. An extension of this from a 3D character turning 90 degrees and being able to explore a 4th Dimension is present in the unreleased game *Miegakure* in which the player can move through the 4th Dimension using the controller, to help orient the player a graphic² is displayed below the controlled character showing the position of the slice of the 4D world, these slices are like the 2D slices we can see in MRI imaging but in 3D.

Whilst these worlds explore pre-made or procedurally generated assets (models, textures, sprites, etc.), and this program instead will explore moving through various parameters, ideas for how the user can interface with the program to understand where exactly they are.

2.5.1 Graphical Prompts

A simple idea to allow the user to recall where in the parameter space they were could be a spider / radar plot, or parallel coordinates plot. For our intents these are the same as we're only displaying a single set of parameters the spider plot is just a 'round' version of the parallel coordinates plot. This could exist on the screen somewhere and change as the user moved around the parameter space, and would allow the user to recall approximately where they were. The spider diagram could also be 'extruded' from how it looked at every point to create a 3D model that represented how you moved through the parameters.

2.5.2 Controls

The above mentioned video games allow for the user to move through the space by only letting the user worry about at most 3 dimensions at any one time, and regarding the others as static when moving through them, this allows the user to control the movement with traditional controls (either a games controller or keyboard and mouse).

²Which the developer tells me is based on an astrolabe

In Figure 2.5, the WASD style of control, popular in video games is used. This will be familiar to people who have played video games but not to people who haven't, for whom it may make sense to use the arrow keys. It's also important to note that in that particular demo the 'd' key increases x , this gives the feeling that the world is moving 'beneath' you instead of you moving the world itself; if the controls were flipped such that 'd' decreased x , it would instead feel like you were moving the world.

Another option would be something like an array of knobs, these could encode one parameter each, and like a scientific instrument be 'tuned'. However this relies on specialist hardware. This would probably feel less like 'moving' through a space and more exploring a range of possibilities.

Chapter 3

Music

3.1 Navigating Sonically

I have developed a theoretical system for using sound to navigate a higher-dimensional space, using ideas from *Harry Partch* and *Joe Monzo's* approached to tuning theory.

This is a process to generate intervals based on a set of spacial coordinates, or simply a set of parameters. Taking inspiration from the *just-intonation* tuning method a ratio comprised of integers only sounds consonant.

Tuning theory is a fundamental question when it comes to developing musical systems, the common myth as told by Iamblichus is that Pythagoras was “walking near a brazier’s shop, he heard from a certain divine casualty the hammers beating out a piece of iron on an anvil, and producing sounds that accorded with each other.” [16, p.62] Pythagoras had discovered that hammers of certain weights were consonant with each other musically. This leads to the concept of intervals.

3.1.1 Intervals

An *interval* is simply a ratio of two frequencies: $\frac{f_1}{f_2}$. Pythagoras’ hammers were in consonant as their weights were 12, 9, 8, and 6. The ratios here of 12 : 6 correspond to $\frac{2}{1}$ or an octave, similarly $12 : 9 = \frac{4}{3}$, $12 : 8 = \frac{3}{2}$ for the other two. Whilst this myth is probably apocryphal, the use of ratios of strings for music have been used since antiquity. What Pythagoras proved is that intervals are these ratios and that materials have properties from which these frequencies arise.

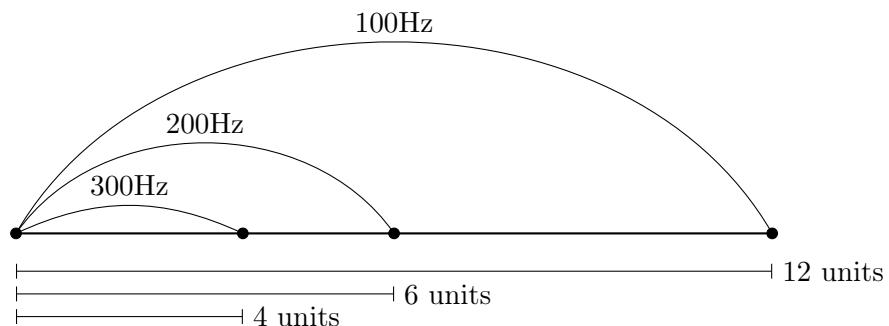


Figure 3.1: Diagram showing intervals being related to the idea of a ratio of lengths for example the 300Hz frequency is related to the 200Hz frequency by $\frac{6}{4} = \frac{3}{2}$

Pythagoras then set about creating a tuning system based around the $\frac{3}{2}$ ratio (the perfect fifth). The perfect fifth is the most consonant ratio that is not unison or an octave¹. Using this ratio we can combine the numbers 2 and 3 to different powers to generate intervals. For example $2^0 3^0 = 1$ or unity, other intervals can be generated but we run into a problem. Using this method of moving in $\frac{3}{2}$ increments if we try to go 7 octaves i.e. $(\frac{2}{1})^7 = 128$ vs 12 fifths² we get $(\frac{3}{2})^{12} = 129.7463379\dots$ the ratio of these two different values is $\frac{(1.5)^{12}}{2^7} \approx 1.01364$. In essence the comma's existence can be surmised by: there exists no $n \neq 0$ such that $2^n = 3^n$. This is important as it marks the difference from the system we use in the west today called '12-tone equal temperament' which divides an octave into 12 equal tones of $2^{\frac{1}{12}}$ to solve this problem.

For our purposes we will be using intervals, ratios of frequencies, to generate new frequencies or notes for use in the program. If we generate a ratio then we may simply multiply a frequency to find the frequency of the other note $x \cdot \frac{f_1}{f_2} = y$.

A p -limit tuning system is a just-intonation based system where every interval's highest prime factor is p . [8, p.76, 109] Pythagoras' tuning system can be said to be 3-limit. These intervals can be represented using an exponent vector, usually called a 'monzo' [7]; for example $3 : 2$ is represented as such: $| -1 1 \rangle$. This is simply a shorthand for $2^{-1} 3^1$, and can be extended: $|e_1 e_2 \dots e_n\rangle$ where each e_i in the vector is an exponent of a prime number $2^{e_1} 3^{e_2} \dots p_n^{e_n}$.

Often these vectors are shown as $| * e_2 \dots e_n\rangle$ The * represents the idea of octave equivalence, a musical idea that a doubling of the frequency produces a note that is 'the same', and as such any value of exponent for the 2 can be placed there. As an example $|1 1\rangle = 2^0 3^1 = \frac{3}{1}$ but this is equivalent to $\frac{3}{2}$, as you can simply halve the frequency produced by the ratio to get back to it. To normalise the interval the same octave, set the constraints $1 \leq \frac{f_1}{f_2} < 2$ and similarly for other octaves higher or lower, just halving and doubling the bounds.

3.1.2 Navigating Space

In 2D

Assume a base frequency, 440Hz = f_1 . As an example, if the user was at $x = 1$, $y = -1$.

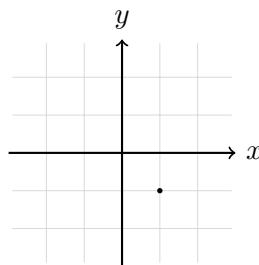


Figure 3.2

¹It is named so because of musical convention, it being five notes between notes in a scale, it has no direct relation to the ratio of 1.5

²These should be the same see: the circle of fifths

This is represented in the vector $|0\ 1\ -1\rangle$ which corresponds to $2^0 3^1 5^{-1} = \frac{6}{5}$. $\frac{6}{5}$ is a minor-third, and represents 528Hz, this is consonant.

Or in general:

$$2^n \cdot 3^x \cdot 5^y \cdot f_1 = f_2$$

If the user is at non-integer numbers for x and y we can calculate the frequencies too.

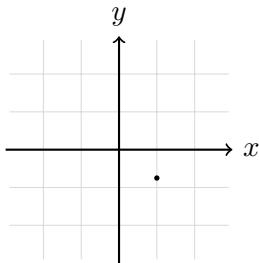


Figure 3.3

Where $x = 1$ and $y = -0.75$ we can simply find the frequency as

$2^0 \cdot 3^1 \cdot 5^{-0.75} \cdot 440\text{Hz} = 394.772\text{Hz}$. This will sound dissonant and not particularly nice.

The idea here is simple, generating intervals like this allows for some points in space to be consonant and others to be dissonant, these could correspond visually with some output being produced by the program but should hopefully allow a user to differentiate between points in space using their sense of hearing.

In Higher Dimensions

This idea extends into higher dimensions naturally, simply adding to the number of terms in the vector $| * e_3 e_5 \cdots e_p \rangle$. Each of these exponents could represent some parameter added to the program.

3.1.3 Limitations

As more variables are introduced integer ratios will sound less ‘strongly consonant’ and may be harder to understand naturally. However, every interval that is possible with less variables will still be possible so a user could still explore one or two parameters at a time and get the same results.

3.1.4 Processing Demo

I have created a demo of this concept in 2D in processing, the code is in Appendix B. The program simply has a integer-marked grid that the user can navigate using WASD and will calculate the interval based on the coordinates of the point. Similarly to Figure 3.3:

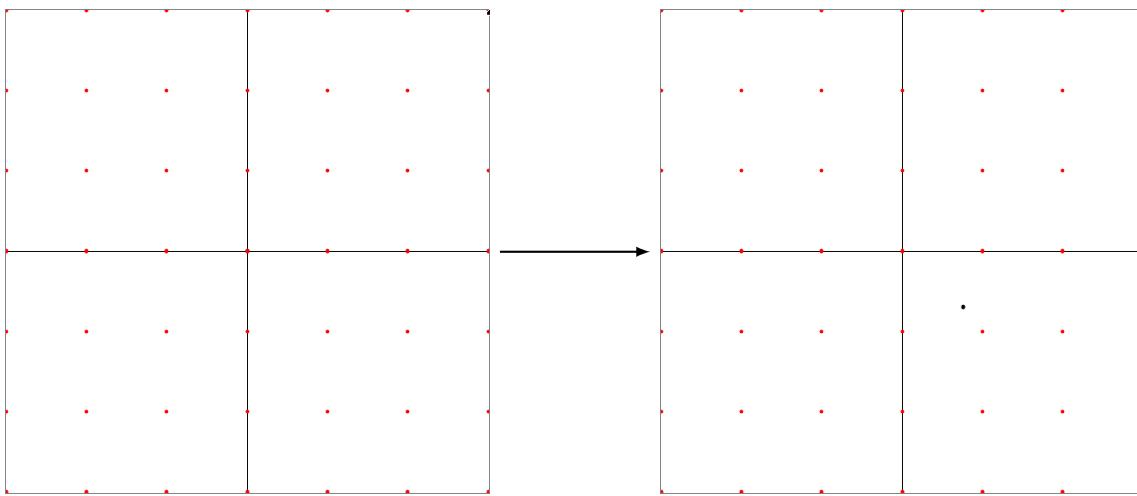


Figure 3.4: At first the black dot is hidden beneath the centre red dot, the user uses WASD to navigate, the audio is out of tune in the second state

This demo is very simple and only uses two sine waves so is very unpleasant to listen to, but illustrates the point, at points outside those marked there is a ‘beating’ sound to the waves as they do not harmonise correctly. However, there are points that aren’t marked that *do* sound somewhat consonant, so this cannot be the only method for navigation, only a prompt to help the user.

3.2 Composition

To create a sound accompaniment to the visual aspect we first need to consider what style of music we should explore, that is to say, create at least an idea of composition. Generative music is the technique we’re exploring here, with musical ideas being emergent from a set of parameters, and created from a system that processes them. The definition is fairly vague as with generative art, the main requirement is that a system is setup and creates the music, this doesn’t need to be a computer, but often is.

Generative music is rather recent, with Brian Eno being major figure popularising its use, he would often use the analogy of a Moiré pattern to describe how the programs would work at the time.

The immediate style of music to draw from Viner’s work would probably be that of

minimalism, the repetition and difference across an image is very similar in style to especially the percussive works of Reich, Glass, Riley, and even the drone works of La Monte Young to an extent. Another, less famous example that I feel conveys the feeling well is *Jon Gibson - Cycles* (1977), the cover for the recording is a Moiré pattern, and the work modulates a 7-note pattern that comes into and out of phase with itself.

In fact, Eno refers to his inspiration of generative music to be triggered by hearing *It's Gonna Rain* by Reich, linking generative music as a concept pretty solidly to Minimalism. Reich used the analogy of fabric work and weaving, featuring on the cover of Music for 18 Musicians is a woven piece of fabric; this seems similar to the idea of a grid (given that weaving takes place on a matrix of strings, perhaps the crossing points could be seen as ‘vertices’)

Overall the idea with the composition should be to enhance what is on the screen, if the image is disordered the sound should be too, if it is ordered and regular the sound should follow. If there are some number of vertices in the shapes on the screen there should be some feedback too.

Further this will be combined with the methods mentioned in section 3.1, including the ideas present in the mentioned minimalist composers. This will be an element in the work that is unique from the art and perhaps helps distance itself from being purely a replica of Viner.

3.3 Synthesis

Broadly there are four approaches to digital synthesis, that of the Processed Recording, the Spectral Model, the Physical Model, and the Abstract Algorithm [14].

For this project spectral and physical modelling are out of the scope and would require more specialist audio software frameworks. Processed recording, includes more sample-based audio and manipulation of that, granular synthesis is an example. The ‘abstract algorithm’ methods include things like FM synthesis, which in its most basic form is comprised of a carrier waveform who’s frequency is modulated by another waveform, this can be extended by things like including feedback at various stages of the processing.

On top of these methods, there should also be a consideration to audio effects, processing and p5.js both have sound libraries with built-in audio effects, reverb and delay perhaps being the most important to create the idea of a space in the sound.

Creating these methods in software are fairly straightforward when there is already some digital signal processing in place; a library called **Tone.JS** handles the practicalities of generating sound in JavaScript and works alongside p5.js nicely.

Additive synthesis makes sense for use here there are a lot of controls over parameters, it is easy to implement and can reproduce a number of sounds. As in the composition section, work like ‘Cycles’ is performed on the organ for example. This is an easy sound to replicate in

additive synthesis as you can achieve the effect by playing multiple sine waves at the same time; in fact this is what the earliest pipe-less organs did, even if not in software, they would have series of plates of disk that encoded the waveforms and played them all at once through a mixer to create complex waveforms [3].

3.3.1 Additive Synthesis in Tone.JS

First I have created a bank of oscillators that can be tuned to any frequency and triggered. I can loop through these and use an ‘envelope’ consisting of to control the volume over time of all of them at once, this is simply set as such:

```
const env = new Tone.AmplitudeEnvelope({
    attack: 1,
    decay: 2,
    sustain: 1,
    release: 0.3
});
```

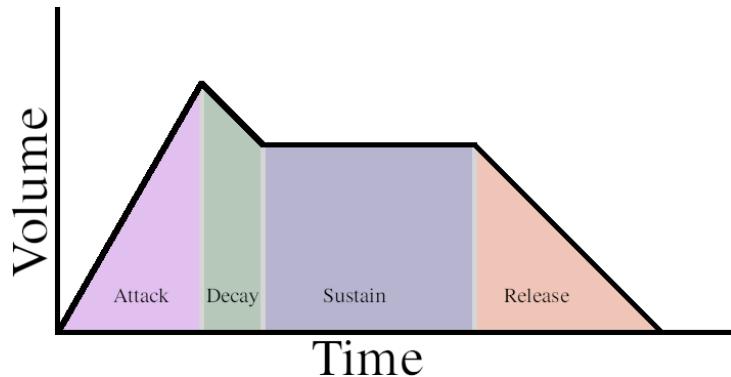


Figure 3.5: Typical attack decay sustain release (ADSR) curve

The number of oscillators in a bank can be chosen at will and for our purposes are tuned as per the harmonic series this is to create a pleasing timbre (the ‘summation’ or additive part is when we feed them into a mixer). Each oscillator is labelled with a number that corresponds to the term of the series they are then when the attack for each oscillator is triggered, the frequency is calculated simply by the incoming note frequency times the harmonic index. Also each oscillator will be quieter than the last according, this is again to create timbre.

Since we are adding these oscillators’ signals together, consideration needs to be given to the phase angle that each lay at. Setting them to `(i / NUM_OSCS) * 360` will allow for there to be no harsh constructive or destructive interference between the oscillators.

3.4 Sequencing

To create a sequence I considered a couple of options, to have a fixed note sequence for example. This would perhaps be too static and not fit with the idea of the project being generative art. Another simple option was to make notes completely random, this tends to sound very ‘robotic’ like a Sci-Fi movie control panel. There could also be a random choice within a set of consonant notes (i.e. that of a chord) this is a better idea because there’s less chance of a ‘bum note’ (one that sounds out of place).

An extension of this idea is to use Markov Chains to model music in [5] for example, they analyse Bach, Mozart, Palestrina, and Beethoven to create a set of transition matrices for their chord progressions. In this example chord progressions are used but the method can easily be adapted to note progressions too.

For my sequence I picked six possible tones, expressed as intervals from a base tone, as 1, 1.2, 1.25, 1.5, 1.6, 2 these represent unison, a minor third, a major third, perfect fifth, and a minor sixth, and the octave (all just-intonation). These tones were picked because they were the intervals in 5-limit that didn’t have recurring decimals.³ Each tone is also then affected by the interval generated above to create a possible series of tones that depend on the parameters of the program.

To create a Markov chain sequence first a matrix of probabilities that represent the chain needs to be created. This defines the possible transitions between notes and the probabilities of them, each row should add to give a total probability of 1. Changing the values in this matrix can therefore be thought of as composing what series of tunes are possible and probable.

$$\begin{array}{ccccccc}
 & 1 & 1.2 & 1.25 & 1.5 & 1.6 & 2 \\
 \begin{matrix} 1 \\ 1.2 \\ 1.25 \\ 1.5 \\ 1.6 \\ 2 \end{matrix} & \left(\begin{array}{cccccc} 0 & 0.2 & 0.5 & 0.1 & 0.1 & 0.1 \\ 0.1 & 0 & 0.4 & 0.1 & 0.2 & 0.2 \\ 0.3 & 0 & 0.1 & 0.4 & 0.1 & 0.1 \\ 0.3 & 0.2 & 0.2 & 0 & 0.1 & 0.2 \\ 0.2 & 0.2 & 0.2 & 0.3 & 0 & 0.1 \\ 0.4 & 0.1 & 0.2 & 0.2 & 0.1 & 0 \end{array} \right)
 \end{array}$$

This represented in code becomes:

```
const markovObject = {
  1: [0 , 0.2, 0.5, 0.1, 0.1, 0.1],
  1.2: [0.1, 0 , 0.4, 0.1, 0.2, 0.2],
  1.25: [0.3, 0 , 0.1, 0.4, 0.1, 0.1],
  1.5: [0.3, 0.2, 0.2, 0 , 0.1, 0.2],
  1.6: [0.2, 0.2, 0.2, 0.3, 0 , 0.1],
  2: [0.4, 0.1, 0.2, 0.2, 0.1, 0 ]
}
```

³Ultimately this is arbitrary but ends up fitting the definition of an augmented scale, this type of scale was used more extensively in the 20th Century in Jazz, it was also used in Liszt’s ‘Faust’s Symphony’. Another way of thinking about it is two major triads above a base pitch (e.g. C E G and E Eb Ab).

```

1.6: [0.2, 0.2, 0.2, 0.3, 0 , 0.1],
2:   [0.4, 0.1, 0.2, 0.2, 0.1, 0 ]
};


```

Using an object allows for fast lookup, then each index in the object refers to an index of the keys of the object. So picking the next is simple, generate a random number from 0 to 1 and then accumulate probability until we reach it:

```

probList.forEach((e, i) => {
  if (picked >= acc && picked < acc + e) {
    intervalIndex = i;
  }
  acc += e;
});

```

Where `probList` is the array referenced by the key of the current note in the object and `intervalIndex` is the index of the array of keys in the object.

Overall this is a very simple method for creating note sequences but is quite effective for limiting the number of possible transitions with a lot of flexibility.

3.5 Chords

For the music to not sound one-dimensional we need to pick more than one note to play in the additive synthesis. This leads to the idea of chords, again we can build these with intervals, this time using a bit more musical theory.

I chose that for each vertex that is present in the middle of the screen there should be another note in the chord. The structure of the chord I chose is simple with the notes being the base note, the minor third, the fifth, the minor seventh, the octave, and the whole tone (or 2nd). Represented by the ratios $1, \frac{6}{5}, \frac{3}{2}, \frac{9}{5}, 2, \frac{9}{8}$ respectively. This is purely an artistic choice. For example if a triangle is in the centre of the screen the first three intervals in the chord will be played $1, \frac{6}{5}, \frac{3}{2}$ resulting in a minor fifth chord. If the user then moves on top of a square the $\frac{9}{5}$ ratio is added to the chord making a minor seventh, similarly if there's a pentagon it simply adds an octave and for a hexagon a whole tone which will sound strange but hexagons are rarer to find, making it feel 'special'.

To implement this I simply store the chord in a function expression that returns the frequency to set each voice (oscillator bank) to in an array up to the index `i`. Then another function expression that triggers the release on voice and attack with the new frequency.

```

const getChord = (i) => [
  base*currentNote,
  base*currentNote*1.2,
  base*currentNote*1.5,

```

```

base*currentNote*1.8,
base*currentNote*2,
base*currentNote*1.125
];

const playVoice = (note, time) => {
  voiceIndex++;
  voiceIndex = voiceIndex % bvoices.length;
  bvoices[voiceIndex].triggerRelease(time);
  mvoices[voiceIndex].triggerRelease(time);
  bvoices[voiceIndex].triggerAttack(note, time);
  mvoices[voiceIndex].triggerAttack(note*interval, time);
};

```

Since the number of possible notes to play at once is 6, the voices array contains six oscillator banks, each with 4 oscillators in them leading to 24 sine waves oscillators at once at maximum. There are two arrays of voices for the base and movable voice so overall there are 48 sine oscillators.

Extensions of this method may lend themselves to being able to control the waveform types and phase of each of the waves as well as what harmonics are present in each voice. Also the ability to change the intervals in a chord. These seem less important for a application such as this primarily static artistic use.

Chapter 4

Recall

4.1 Saving the State

To be able to recall past states a system needs to be in place that saves parameters periodically. Every 100 frames a new ‘snapshot’ of the state is generated; this is an element in a JSON tree. Given the program is running at around 60Hz (it is, however, not fixed) this means that approximately every 1.6 seconds the state is saved.

However, if we left the program running memory use would grow, so when saving the state the program should check to see if it’s changed or not. This has the effect of compressing time in a neat way because you can interact with the program, then step away and come back, begin interaction again and still be able to use the recall as if there was no large gap between times.

The naive implementation of this is to create a json object that looks something like this:

```
{  
    "seed": 230128038,  
    "uuid": 102830128,  
    "elements": [  
        {  
            "x": 28,  
            "y": 0,  
            "z": 0,  
            "noiseLevel": 0.4,  
            "gridSpacing": 0.3  
        },  
        {  
            "x": 31,  
            "y": 10,  
            "z": 0,  
            "noiseLevel": 0.4,  
            "gridSpacing": 0.3  
        }  
    ]  
}
```

But, this approach of only using an array doesn’t work, as when you ‘rewind’ you’re actually branching from the progression that is ahead in the array. If you only insert at the end of the

array then this history gets mixed up leading to a mess of non-chronological states.

So to fix this, a method using a tree should be devised. Each time you ‘rewind’ the tree should create a new branch from which you can begin control and generate state elements in that branch. To do this I used the library **Data-Tree**. This helps manage the tree structure by using references, and can export and import the structure to and from JSON. The library also provides methods to search and traverse the tree, operating on each node.

Exporting the tree to a file is then very simple, create a method that exports the JSON tree and downloads it to the computer, the filename can be used to store the seed. This works by creating a blob (binary large object), url to the blob and then a link element that downloads the json file, then using JavaScript to click the link and finally removing the link.

4.2 Drawing and Interaction

Drawing the tree to the screen requires traversal. If we use depth-first traversal we can work out where a new branch is formed by checking if the depth of the node is the previous node’s depth plus one. Then we can simply move down and draw a point. If not, we branch, however it becomes more complicated; at every node we store the (x, y) values into a json object with the key being the node’s unique key. Then when we see that the tree has branched we can find the parent node (**Data-Tree** stores this) and set the y value to this and the x values increases by two times the diameter of the point. We then also need to draw a line from the new branch’s first node to the point where it branches, this can be done by simply recalling the x value of the parent node. Finally, the most recent node is coloured red. In the end it looks like:

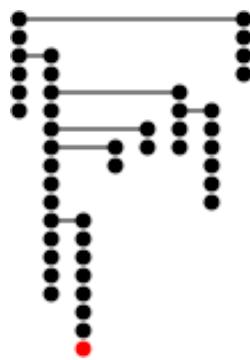


Figure 4.1

To interact with the tree, each point is modelled as a button. These buttons have a `display()` method that checks if the mouse is hovering and then draws a point. The point becomes twice the size if hovered and uses the key of the node to check whether to make itself red or not. The button can be used to check when the mouse is clicked if the button is hovered at the time of clicking and then the state can be loaded.

4.3 Recalling Parameters

Recalling parameters is simple, when we know the key of the node to recall (from the button that was clicked on) we can search the tree and then return the values stored at that node. The program then sets each of the variables in the main script to these values.

Importing a tree from json is similar, except we use `Data-Tree` to import the file into it's data structure first and then load the state at the last node that was imported. We also call `noiseSeed(filename)` (with the `.json` stripped) to set the seed correctly.

Chapter 5

Ethics

This project does not handle any personal data, nor deals with identification of people from data. There is no subject matter that may offend or otherwise effect groups of people.

However there is a question over ownership and who owns images produced by this program if they are indistinguishable from Viner's own work. One point to consider is that where his work is pen-plotter images on listing paper, this is purely computer generated imagery.

Of course, also the provenance associated with the work is also completely different; in the case of Prado's Mona Lisa for example, it's clear that despite being made around the same time, in the same workshop, and of the same subject it is a separate piece of work with a different story. This perhaps is not the best example however the creator of the Prado version likely worked in da Vinci's studio at the time [13].

It's also worth mentioning the fair dealing exception to UK copyright law which says that you are allowed to "copy limited extracts of work for non-commercial research or private study"[11]. However I am unsure if this legal protection extends to *the potential for generating* copies of work. It is clear to me however that I am not attempting to reproduce directly and instead explore the techniques of creation of this work using completely different technology.

References

- [1] E. Abbott. *Flatland: A Romance of Many Dimensions*. Roberts Brothers, 1885.
- [2] D. Asimov. The grand tour: a tool for viewing multidimensional data. *SIAM journal on scientific and statistical computing*, 6(1):128–143, 1985.
- [3] P. Comerford. Simulating an organ with additive synthesis. *Computer Music Journal*, 17(2):55–65, 1993.
- [4] D. T. L. Deborah F. Swayne, Andreas Buja. Exploratory visual analysis of graphs in GGobi. In *Exploratory Visual Analysis of Graphs in GGobi*, 2003.
- [5] P. Kiefer and M. Riehl. Markov chains of chord progressions. *Ball State Undergraduate Mathematics Exchange*, 10:16–21, 2016.
- [6] R. Lycett. Darrell viner: Materiality, process and the coded object, Jan. 2016. Abstract Retrieved 05/04/2021: <https://eprints.hud.ac.uk/id/eprint/28205/>.
- [7] J. Monzo. Monzo / exponent vector / prime-exponent vector, 2005. <http://tonalsoft.com/monzo/lattices/lattices.htm>.
- [8] H. Partch. *Genesis Of A Music*. Da Capo Press, 1974.
- [9] K. Perlin. Noise and turbulence, 1998. Retrieved 10/03/2021: <https://mrl.cs.nyu.edu/~perlin/doc/oscar.html>.
- [10] K. Perlin. Improving noise. In *Proceedings of the 29th Annual Conference on Computer Graphics and Interactive Techniques*, SIGGRAPH '02, pages 681–682, New York, NY, USA, 2002. Association for Computing Machinery.
- [11] gov.uk. Exceptions to copyright. [#overview](https://www.gov.uk/guidance/exceptions-to-copyright).
- [12] Henry Moore Institute. Darrell viner: Early work, 2011. <https://www.henry-moore.org/whats-on/2011/07/27/darrell-viner-early-work>.
- [13] Museo Nacional del Prado. The Mona Lisa - The Collection - Museo Nacional del Prado. <https://www.museodelprado.es/en/the-collection/art-work/the-mona-lisa/80c9b279-5c80-4d29-b72d-b19cdca6601c>.
- [14] J. O. Smith, III. Viewpoints on the history of digital synthesis, Dec. 2005. Keynote paper from proceedings of international computer music conference pp. 1-10, Oct. 1991. Revised for I'IRCAM, 1992.
- [15] J. Stell. Notes, 2019. Unpublished.

- [16] T. Taylor. *Iamblichus' Life of Pythagoras, or Pythagoric Life*. J.M. Watkins, 1818.
- [17] J. A. Vince. *Picaso A Computer Language For Art And Design*. PhD thesis, Brunel University, 1975. Available: <https://eprints.mdx.ac.uk/10169/>.
- [18] D. Viner. Artist statement. Retrieved 14/12/2017:
<http://www.darrellviner.org/artiststatement/>.

Appendix A

External Materials

- Libraries
 - p5.js javascript library: <https://p5js.org/>
 - tone.js javascript library: <https://tonejs.github.io/>
 - Data-Tree javascript library: <http://cchandurkar.github.io/Data-Tree/>
- Other External Materials
 - Some of the initial graphics code was based on notes by my supervisor, John Stell
 - The code for the additive synth was based on and heavily modified from
<https://github.com/ejarzo/additive-synth>

Appendix B

Audio Demo

```
import processing.sound.*;  
  
SinOsc base;  
SinOsc variable;  
float baseFreq = 440;  
// they are in unity to start with  
float variableFreq = 440;  
  
float volume = 0.2;  
float s_x,s_y,x,y = 0;  
float interval;  
  
boolean isUp, isDown, isLeft, isRight;  
  
void settings() {  
    size(501, 501);  
}  
  
void setup() {  
    base = new SinOsc(this);  
    base.play();  
    base.freq(baseFreq);  
    base.amp(volume);  
  
    variable = new SinOsc(this);  
    variable.play();  
    variable.amp(volume);  
}  
  
void draw() {  
    background(255);  
    translate(width/2, height/2);  
    scale(1, -1);  
  
// make some axes  
    stroke(0);  
    strokeWeight(1);
```

```
line(0,height/2,0,-height/2);
line(width/2,0,-height,0);

strokeWeight(4);
point(x,y);

// mark where the integer values will be, for demo reasons
stroke(255,0,0);
strokeWeight(4);
for(int i = -3; i <= 3; i++) {
    for(int j = -3; j <= 3; j++) {
        point(map(i, -3, 3, -width/2, width/2),
              map(j, -3, 3, -height/2, height/2));
    }
}

// scale x and y to some reasonable numbers -3 -> 3
s_x = map(x, -width/2, width/2, -3, 3);
s_y = map(y, -height/2, height/2, -3, 3);
// now we need to find a way to scale until
// we're in the same octave
interval = pow(3,s_x) * pow(5,s_y);

while (interval < 1 || interval >= 2) {
    if (interval < 1) {
        interval = interval * 2;
    } else if (interval > 2) {
        interval = interval / 2;
    }
}

variableFreq = interval * baseFreq;
variable.freq(variableFreq);

if (isUp) y++;
if (isDown) y--;
if (isRight) x++;
if (isLeft) x--;
}

void keyPressed() {
    setMove(key, true);
}
```

```
void keyReleased() {
    setMove(key, false);
}

void setMove(char k, boolean b) {
    if (key == 'w') isUp = b;
    if (key == 's') isDown = b;
    if (key == 'a') isLeft = b;
    if (key == 'd') isRight = b;
}
```